

# Chapter 9

TRANSACTION MANAGEMENT AND  
CONCURRENCY CONTROL

# In this chapter, you will learn:

- ▶ Read/Write conflict scenario
- ▶ Locking Methods
- ▶ Lock Granularity
- ▶ Lock Types
- ▶ Concurrency Control with Time Stamping Methods
- ▶ Data Recovery Management

# Read/Write Conflict Scenarios: Conflicting Database Operations Matrix

**TABLE 9.9** READ/WRITE CONFLICT SCENARIOS: CONFLICTING DATABASE OPERATIONS MATRIX

Operations	TRANSACTIONS		RESULT
	T1	T2	
	Read	Read	No conflict
	Read	Write	Conflict
	Write	Read	Conflict
	Write	Write	Conflict

# Concurrency Control with Locking Methods

- ▶ Lock
  - ▶ Guarantees exclusive use of a data item to a current transaction
  - ▶ Required to prevent another transaction from reading inconsistent data
- ▶ Lock manager
  - ▶ Responsible for assigning and policing the locks used by the transactions

# Lock Granularity

- ▶ Indicates the level of lock use
- ▶ Locking can take place at the following levels:
  - ▶ Database
  - ▶ Table
  - ▶ Page
  - ▶ Row
  - ▶ Field (attribute)

# Lock Granularity (continued)


- ▶ Database-level lock
  - ▶ Entire database is locked
- ▶ Table-level lock
  - ▶ Entire table is locked
- ▶ Page-level lock
  - ▶ Entire diskpage is locked

# Lock Granularity (continued)

- ▶ Row-level lock
  - ▶ Allows concurrent transactions to access different rows of the same table, even if the rows are located on the same page
- ▶ Field-level lock
  - ▶ Allows concurrent transactions to access the same row, as long as they require the use of different fields (attributes) within that row

# A Database-Level Locking Sequence

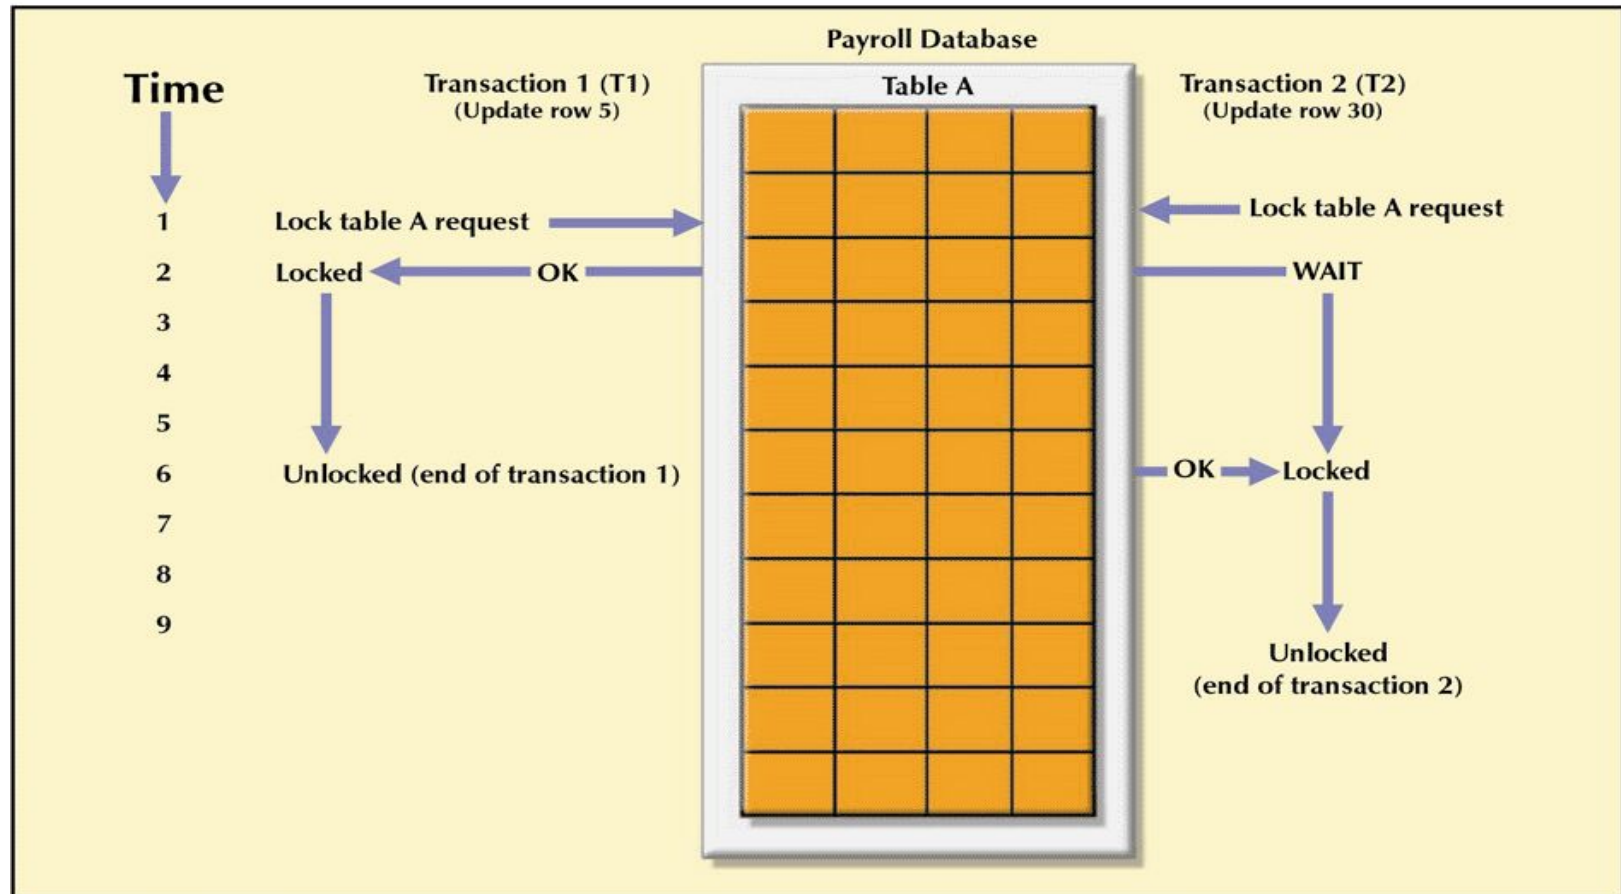
FIGURE 9.3 A DATABASE-LEVEL LOCKING SEQUENCE





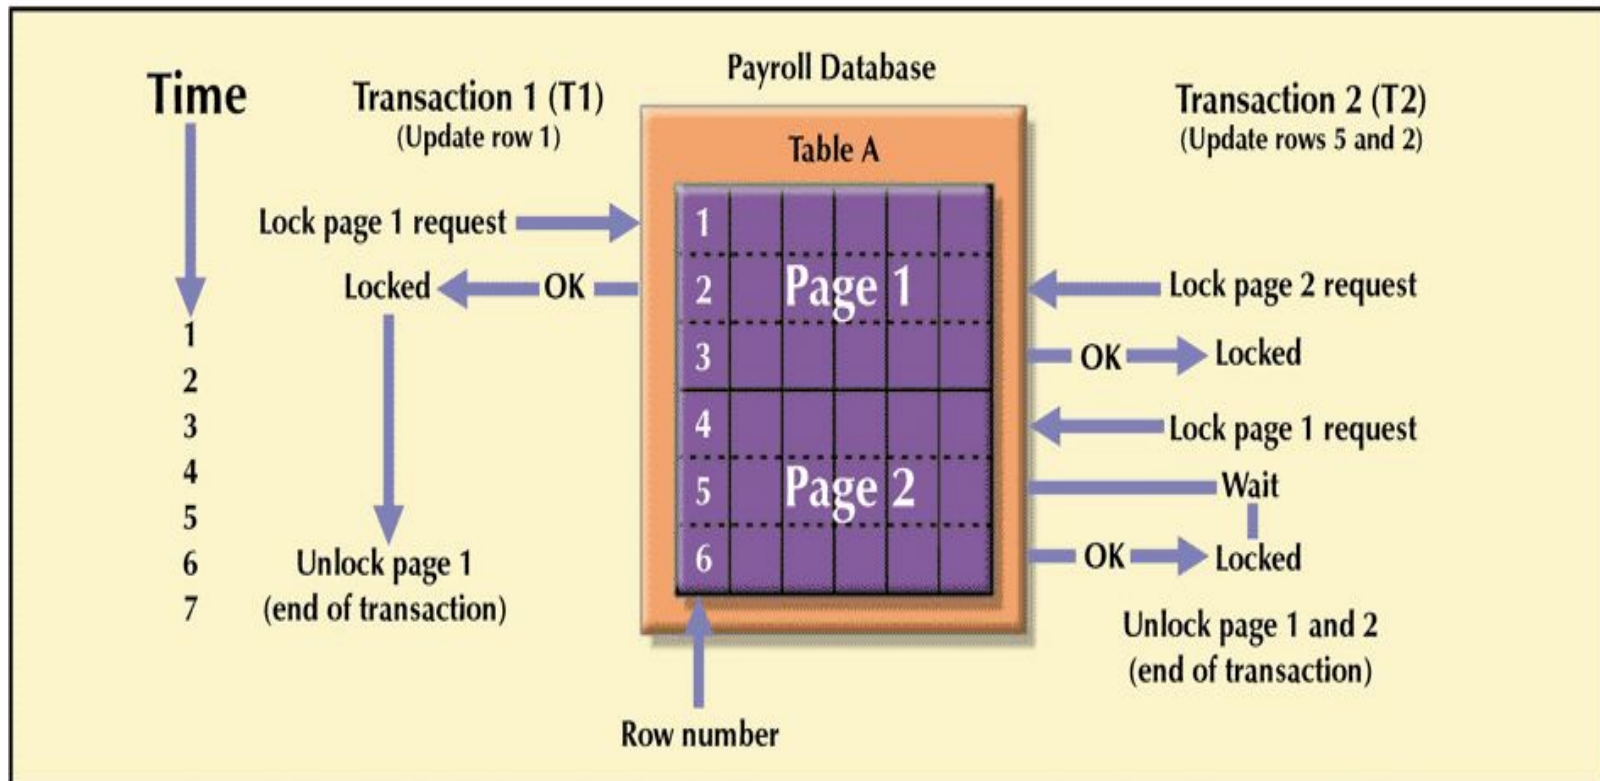
# An Example of a Table-Level Lock

FIGURE 9.4 AN EXAMPLE OF A TABLE-LEVEL LOCK



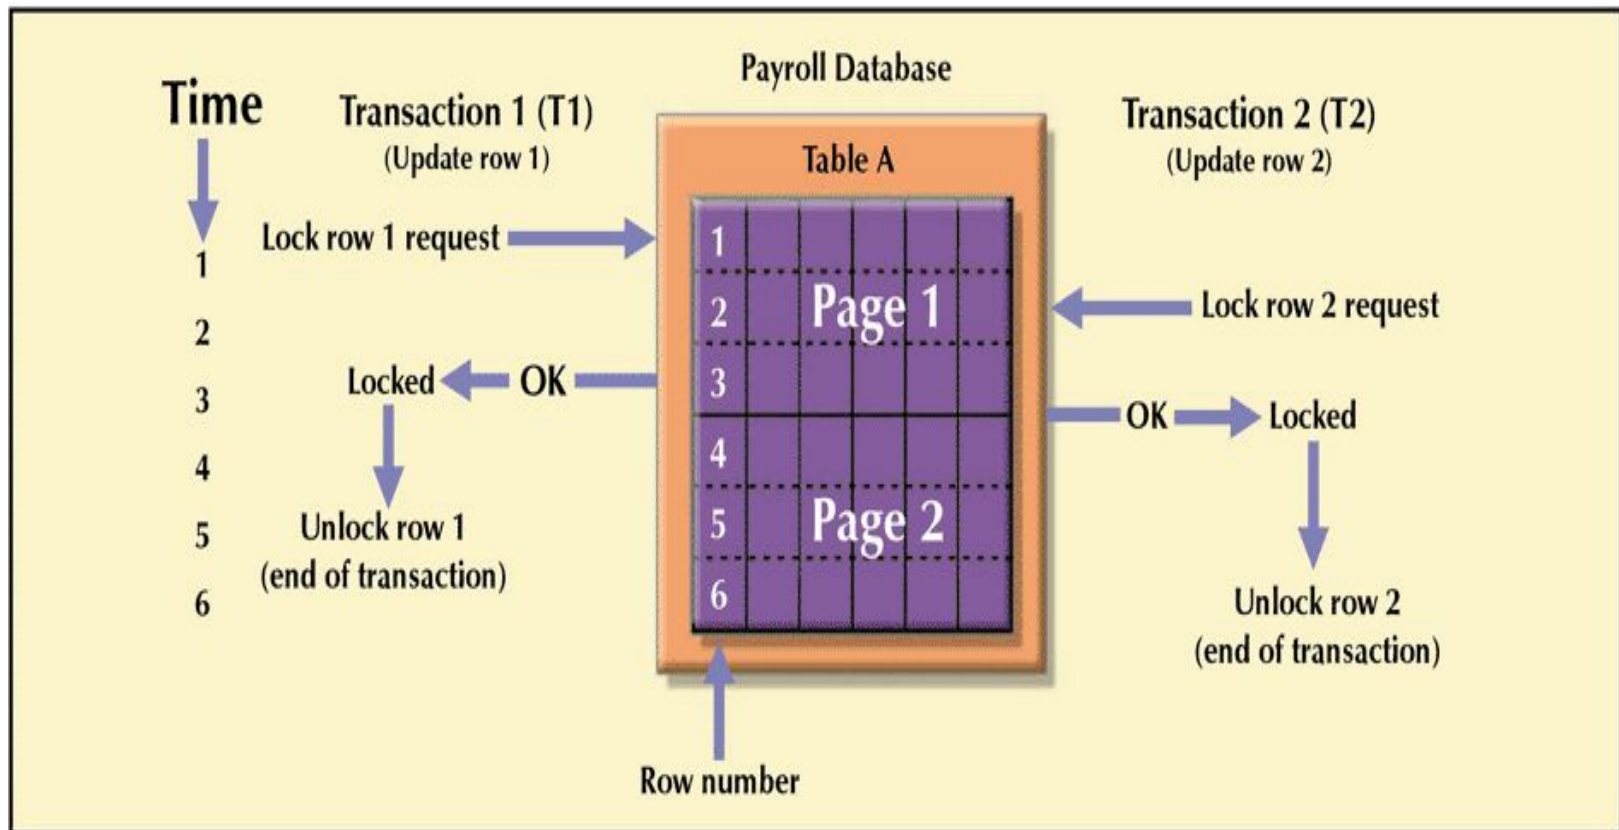
# Example of a Page-Level Lock

FIGURE 9.5 AN EXAMPLE OF A PAGE-LEVEL LOCK



# An Example of a Row-Level Lock

FIGURE 9.6 AN EXAMPLE OF A ROW-LEVEL LOCK



# Lock Types

- ▶ Binary lock
  - ▶ Has only two states: locked (1) or unlocked (0)
- ▶ Exclusive lock
  - ▶ Access is specifically reserved for the transaction that locked the object
  - ▶ Must be used when the potential for conflict exists
- ▶ Shared lock
  - ▶ Concurrent transactions are granted Read access on the basis of a common lock

# An Example of a Binary Lock

**TABLE 9.10** AN EXAMPLE OF A BINARY LOCK

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Lock PRODUCT	
2	T1	Read PROD_QOH	15
3	T1	$\text{PROD\_QOH} = 15 + 10$	
4	T1	Write PROD_QOH	25
5	T1	Unlock PRODUCT	
6	T2	Lock PRODUCT	
7	T2	Read PROD_QOH	23
8	T2	$\text{PROD\_QOH} = 23 - 10$	
9	T2	Write PROD_QOH	13
10	T2	Unlock PRODUCT	

# Two-Phase Locking to Ensure Serializability

- ▶ Defines how transactions acquire and relinquish locks
- ▶ Guarantees serializability, but it does not prevent deadlocks
  - ▶ Growing phase, in which a transaction acquires all the required locks without unlocking any data
  - ▶ Shrinking phase, in which a transaction releases all locks and cannot obtain any new lock

# Two-Phase Locking to Ensure Serializability (continued)

- ▶ Governed by the following rules:
  - ▶ Two transactions cannot have conflicting locks
  - ▶ No unlock operation can precede a lock operation in the same transaction
  - ▶ No data are affected until all locks are obtained—that is, until the transaction is in its locked point

# Two-Phase Locking Example

(a)

$T_1$	$T_2$
<code>read_lock(Y);</code> <code>read_item(Y);</code> <code>unlock(Y);</code> <code>write_lock(X);</code> <code>read_item(X);</code> <code>X := X + Y;</code> <code>write_item(X);</code> <code>unlock(X);</code>	<code>read_lock(X);</code> <code>read_item(X);</code> <code>unlock(X);</code> <code>write_lock(Y);</code> <code>read_item(Y);</code> <code>Y := X + Y;</code> <code>write_item(Y);</code> <code>unlock(Y);</code>

(b)

Initial values:  $X=20, Y=30$

Result serial schedule  $T_1$   
followed by  $T_2$ :  $X=50, Y=80$

Result of serial schedule  $T_2$   
followed by  $T_1$ :  $X=70, Y=50$

- Transaction that support 2pl

$T_1'$	$T_2'$
<code>read_lock(Y);</code> <code>read_item(Y);</code> <code>write_lock(X);</code> <code>unlock(Y)</code> <code>read_item(X);</code> <code>X := X + Y;</code> <code>write_item(X);</code> <code>unlock(X);</code>	<code>read_lock(X);</code> <code>read_item(X);</code> <code>write_lock(Y);</code> <code>unlock(X)</code> <code>read_item(Y);</code> <code>Y := X + Y;</code> <code>write_item(Y);</code> <code>unlock(Y);</code>



# Two-Phase Locking to Ensure Serializability (continued)

Suppose two-phase locking does not ensure serializability. Then there exists a set of transactions  $T_0, T_1 \dots T_{n-1}$  which obey 2PL and which produce a nonserializable schedule. A non-serializable schedule implies a cycle in the

precedence graph, and we shall show that 2PL cannot produce such cycles.

Without loss of generality, assume the following cycle exists in the precedence

graph:  $T_0 \rightarrow T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_{n-1} \rightarrow T_0$ . Let  $\alpha_i$  be the time at which  $T_i$  obtains

its last lock (i.e.  $T_i$ 's lock point). Then for all transactions such that  $T_i \rightarrow T_j$ ,


$\alpha_i < \alpha_j$ . Then for the cycle we have

$\alpha_0 < \alpha_1 < \alpha_2 < \dots < \alpha_{n-1} < \alpha_0$

Since  $\alpha_0 < \alpha_0$  is a contradiction, no such cycle can exist. Hence 2PL cannot produce non-serializable schedules. Because of the property that for all transactions such that  $T_i \rightarrow T_j$ ,  $\alpha_i < \alpha_j$ , the lock point ordering of the transactions is also a topological sort ordering of the precedence graph. Thus transactions can be serialized according to their lock points.

# Two-Phase Locking Protocol

FIGURE 9.7 TWO-PHASE LOCKING PROTOCOL



# Deadlocks

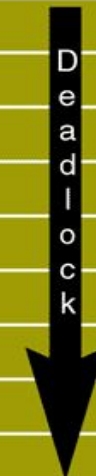
- ▶ Condition that occurs when two transactions wait for each other to unlock data
- ▶ Possible only if one of the transactions wants to obtain an exclusive lock on a data item
  - ▶ No deadlock condition can exist among *shared* locks
- ▶ Control through
  - ▶ Prevention
  - ▶ Detection
  - ▶ Avoidance

# How a Deadlock Condition Is Created

**TABLE 9.11** HOW A DEADLOCK CONDITION IS CREATED

TIME	TRANSACTION	REPLY	LOCK STATUS	
			Data X	Data Y
0			Unlocked	Unlocked
1	T1:LOCK(X)	OK	Unlocked	Unlocked
2	T2: LOCK(Y)	OK	Locked	Unlocked
3	T1:LOCK(Y)	WAIT	Locked	Locked
4	T2:LOCK(X)	WAIT	Locked	Locked
5	T1:LOCK(Y)	WAIT	Locked	Locked
6	T2:LOCK(X)	WAIT	Locked	Locked
7	T1:LOCK(Y)	WAIT	Locked	Locked
8	T2:LOCK(X)	WAIT	Locked	Locked
9	T1:LOCK(Y)	WAIT	Locked	Locked
...	.....	.....	.....	.....
...	.....	.....	.....	.....
...	.....	.....	.....	.....
...	.....	.....	.....	.....

Deadlock



# Concurrency Control with Time Stamping Methods

- ▶ Assigns a global unique time stamp to each transaction
- ▶ Produces an explicit order in which transactions are submitted to the DBMS
- ▶ Uniqueness
  - ▶ Ensures that no equal time stamp values can exist
- ▶ Monotonicity
  - ▶ Ensures that time stamp values always increase

# Deadlock Prevention Techniques

- ▶ **Wait-die scheme:** It is a non-preemptive technique for deadlock prevention. When transaction  $T_i$  requests a data item currently held by  $T_j$ ,  $T_i$  is allowed to wait only if it has a timestamp smaller than that of  $T_j$ . (That is  $T_i$  is older than  $T_j$ ), otherwise  $T_i$  is rolled back (dies)
- ▶ **For example:**
- ▶ Suppose that transaction  $T_{22}$ ,  $T_{23}$ ,  $T_{24}$  have time-stamps 5, 10 and 15 respectively. If  $T_{22}$  requests a data item held by  $T_{23}$  then  $T_{22}$  will wait. If  $T_{24}$  requests a data item held by  $T_{23}$ , then  $T_{24}$  will be rolled back.

# Wait/Die and Wound/Wait Schemes

- ▶ **Wound-wait scheme:** It is a preemptive technique for deadlock prevention. It is a counterpart to the wait-die scheme. When Transaction  $T_i$  requests a data item currently held by  $T_j$ ,  $T_i$  is allowed to wait only if it has a timestamp larger than that of  $T_j$ , otherwise  $T_j$  is rolled back ( $T_j$  is wounded by  $T_i$ )
- ▶ **For example:**
- ▶ Suppose that Transactions  $T_{22}$ ,  $T_{23}$ ,  $T_{24}$  have time-stamps 5, 10 and 15 respectively. If  $T_{22}$  requests a data item held by  $T_{23}$ , then data item will be preempted from  $T_{23}$  and  $T_{23}$  will be rolled back. If  $T_{24}$  requests a data item held by  $T_{23}$ , then  $T_{24}$  will wait.

# Wait/Die and Wound/Wait Concurrency Control Schemes

**TABLE 9.12** WAIT/DIE AND WOUND/WAIT CONCURRENCY CONTROL SCHEMES

TRANSACTION REQUESTING LOCK	TRANSACTION OWNING LOCK	WAIT/DIE SCHEME	WOUND/WAIT SCHEME
T1 (11548789)	T2 (19562545)	<ul style="list-style-type: none"><li>• T1 waits until T2 is completed and T2 releases its locks.</li></ul>	<ul style="list-style-type: none"><li>• T1 preempts (rolls back) T2. T2 is rescheduled using the same time stamp.</li></ul>
T2 (19562545)	T1 (11548789)	<ul style="list-style-type: none"><li>• T2 Dies (rolls back).</li><li>• T2 is rescheduled using the same time stamp.</li></ul>	<ul style="list-style-type: none"><li>• T2 waits until T1 is completed and T1 releases its locks.</li></ul>



# Concurrency Control with Optimistic Methods

- ▶ Optimistic approach
  - ▶ Based on the assumption that the majority of database operations do not conflict
  - ▶ Does not require locking or time stamping techniques
  - ▶ Transaction is executed without restrictions until it is committed
  - ▶ Phases are read, validation, and write

# Database Recovery Management

- ▶ Database recovery
  - ▶ Restores database from a given state, usually inconsistent, to a previously consistent state
  - ▶ Based on the atomic transaction property
    - ▶ All portions of the transaction must be treated as a single logical unit of work, in which all operations must be applied and completed to produce a consistent database
  - ▶ If transaction operation cannot be completed, transaction must be aborted, and any changes to the database must be rolled back (undone)

# Transaction Recovery

- ▶ Makes use of deferred-write and write-through
- ▶ Deferred write
  - ▶ Transaction operations do not immediately update the physical database
  - ▶ Only the transaction log is updated
  - ▶ Database is physically updated only after the transaction reaches its commit point using the transaction log information

# Transaction Recovery (continued)

- ▶ Write-through
  - ▶ Database is immediately updated by transaction operations during the transaction's execution, even before the transaction reaches its commit point

# A Transaction Log for Transaction Recovery Examples

**TABLE 9.13 A TRANSACTION LOG FOR TRANSACTION RECOVERY EXAMPLES**

TRL ID	TRX NUM	PREV PTR	NEXT PTR	OPERATION	TABLE	ROW ID	ATTRIBUTE	BEFORE VALUE	AFTER VALUE
341	101	Null	352	START	****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	****Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, ...
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, ...
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423	CHECKPOINT								
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, ...
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	****Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									

# Summary

- ▶ Transaction
  - ▶ Sequence of database operations that access the database
  - ▶ Represents real-world events
  - ▶ Must be a logical unit of work
    - ▶ No portion of the transaction can exist by itself
  - ▶ Takes a database from one consistent state to another
    - ▶ One in which all data integrity constraints are satisfied

## Summary (continued)

- ▶ SQL provides support for transactions through the use of two statements: COMMIT and ROLLBACK
- ▶ Concurrency control coordinates the simultaneous execution of transactions
- ▶ Scheduler is responsible for establishing order in which concurrent transaction operations are executed

## Summary (continued)

- ▶ Lock guarantees unique access to a data item by a transaction
- ▶ Database recovery restores the database from a given state to a previous consistent state



# Summary (continued)

- ▶ Dirty read:
- ▶ Nonrepeatable reads
- ▶ Phantom Reads


## Isolation Levels: Summary

Trans:

	dirty reads	nonrepeatable reads	phantoms
Read Uncommitted	Y	Y	Y
Read Committed	N	Y	Y
Repeatable Read	N	N	Y
Serializable	N	N	N

weak

strong

- 
- ▶ Phantom read: This means that if you execute a query at time T1 and re-execute it at time T2,
  - ▶ additional rows may have been added to the database, which will affect your results. This
  - ▶ differs from the nonrepeatable read in that with a phantom read, data you already read has not
  - ▶ been changed, but rather that more data satisfies your query criteria than before.



Queries ?